Engenharia de Computação

Especialização Lato Sensu em Ciência de Dados e Analytics

Soluções em Processamento para Big Data

{ Prática Google Data Proc }

Prof. Jairson Rodrigues
jairson.rodrigues@univasf.edu.br

{ google data proc }

AGENDA

Cluster Google

{ roteiro }

- Cadastro no Google DataProc
- Construir um cluster de quatro máquinas
- Acessar o Name Node
- Executar algoritmos de ML
  - ETL
  - Regressão Logística
  - Árvores de Decisão
  - Kmeans

{ google data proc }

https://cloud.google.com/dataproc/

CLOUD DATAPROC

A faster, easier, more cost-effective way to run Spark and Hadoop

VIEW CLOUD DATAPROC DOCS

VIEW MY CONSOLE

Cloud-native Hadoop & Spark

Cloud Dataproc is a fast, easy-to-use, fully-managed cloud service for running **Apache Spark** and **Apache Hadoop** clusters in a simpler, more cost-efficient way. Operations that used to take hours or days take seconds or minutes instead, and you pay only for the resources you use (with per-second billing). Cloud Dataproc also easily integrates with other Google Cloud Platform (GCP) services, giving you a powerful and complete platform for data processing, analytics and machine learning.

{ google data proc }

https://console.cloud.google.com/dataproc/

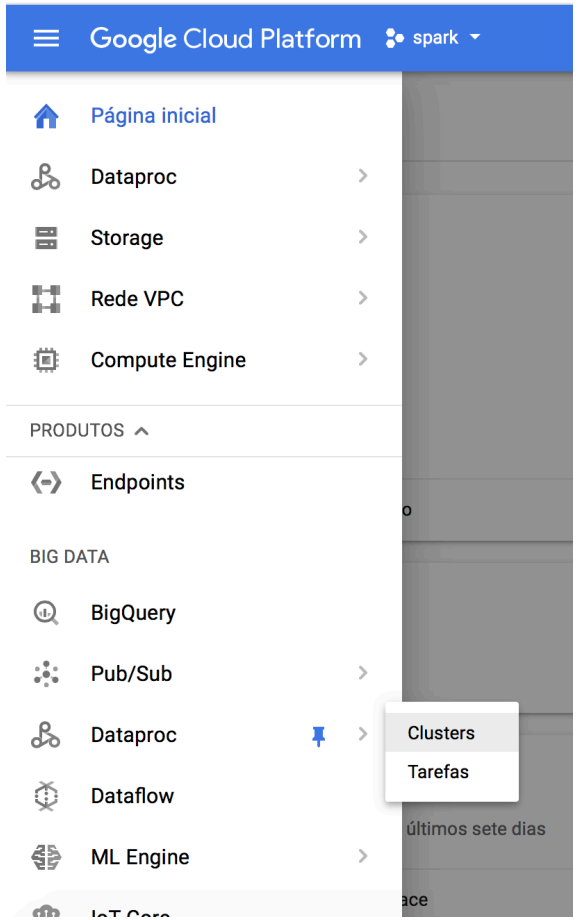

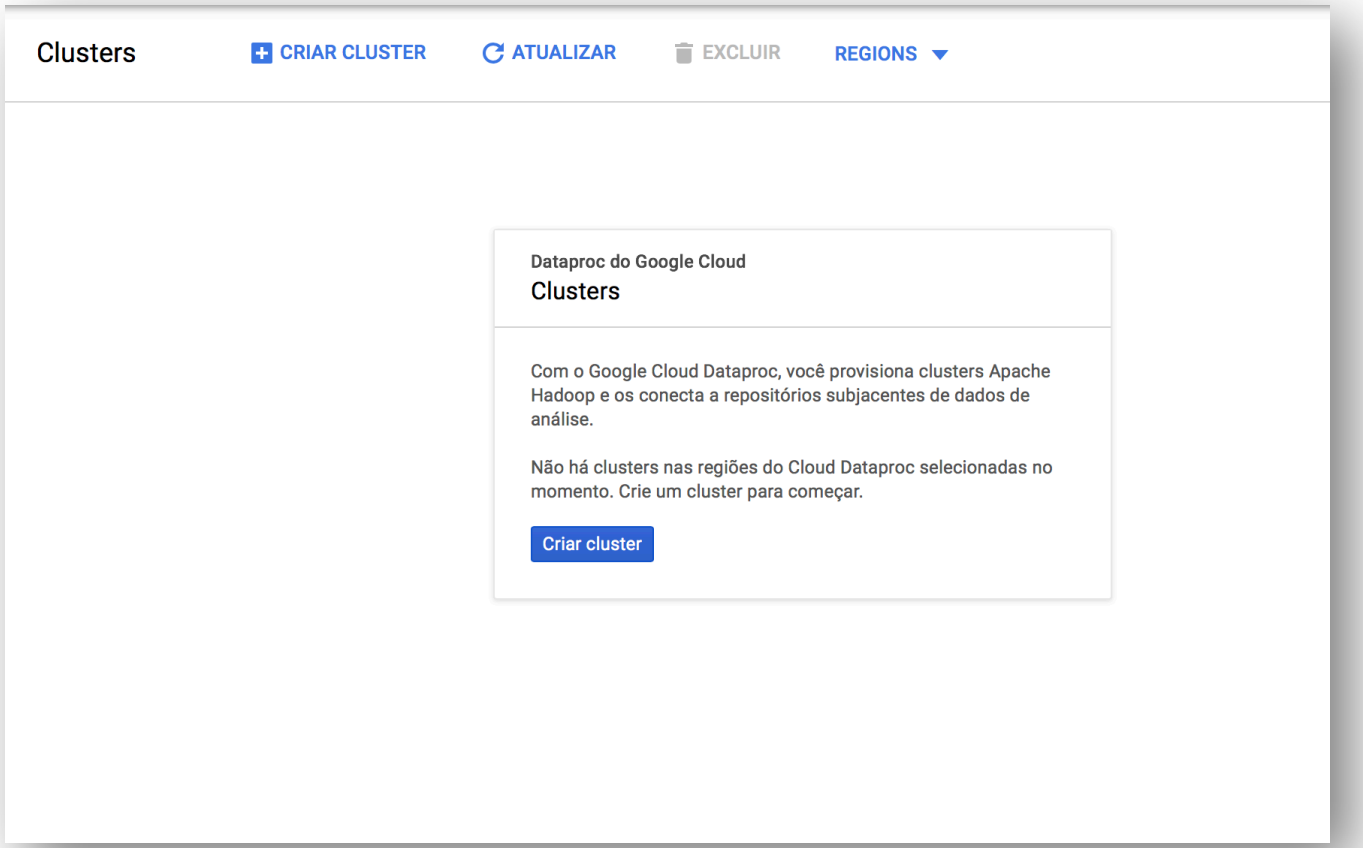

{ google data proc }

Tipos de Máquina Virtual

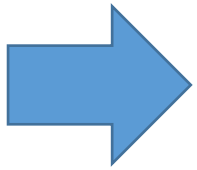

- n1-standard-1 (1 vCPU, 3.75 GB de memória)
- n1-standard-2 (2 vCPU, 7.50 GB de memória)
- n1-standard-4 (4 vCPU, 15.0 GB de memória)
- n1-standard-8 (8 vCPU, 30.0 GB de memória)
- n1-standard-16 (16 vCPU, 60.0 GB de memória)
- n1-standard-32 (32 vCPU, 120 GB de memória)
- n1-standard-64 (64 vCPU, 240 GB de memória)
- n1-highcpu-4 (4 vCPU, 3.60 GB de memória)
- n1-highcpu-8 (8 vCPU, 7.20 GB de memória)
- n1-highcpu-16 (16 vCPU, 14.4 GB de memória)
- n1-highcpu-32 (32 vCPU, 28.8 GB de memória)
- n1-highcpu-64 (64 vCPU, 57.6 GB de memória)
- n1-highmem-2 (2 vCPU, 13.0 GB de memória)
- n1-highmem-4 (4 vCPU, 26.0 GB de memória)
- n1-highmem-8 (8 vCPU, 52.0 GB de memória)
- n1-highmem-16 (16 vCPU, 104 GB de memória)
- n1-highmem-32 (32 vCPU, 208 GB de memória)
- n1-highmem-64 (64 vCPU, 416 GB de memória)

{ configuração de nós }

- Master: Google VM Machine
  - n1-highmem-2
  - CPU: 2
  - RAM: 13 Gb RAM
  - Disco local: 500 Gb
  - # instâncias: 1
- Slave: Google VM Machine
  - n1-highmem-2
  - CPU: 2
  - RAM: 13 Gb RAM
  - Disco local: 500 Gb
  - Disco SSD 375 Gb
  - # instâncias: 3

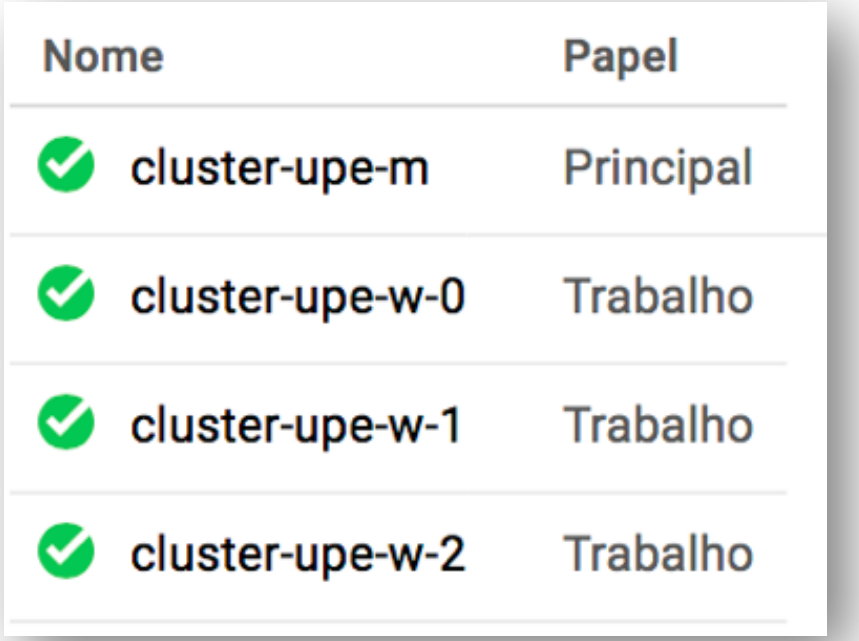

{ google data proc – master }

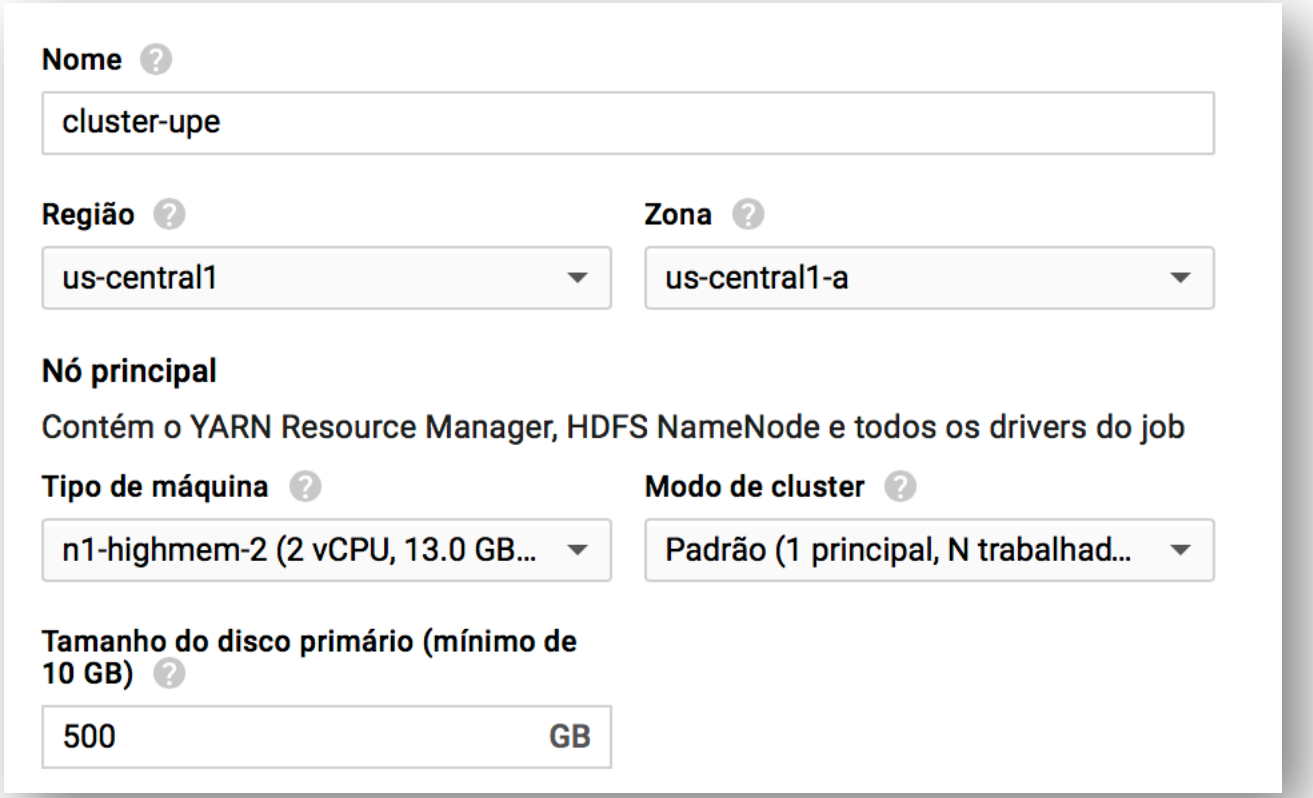

{ google data proc – slaves }

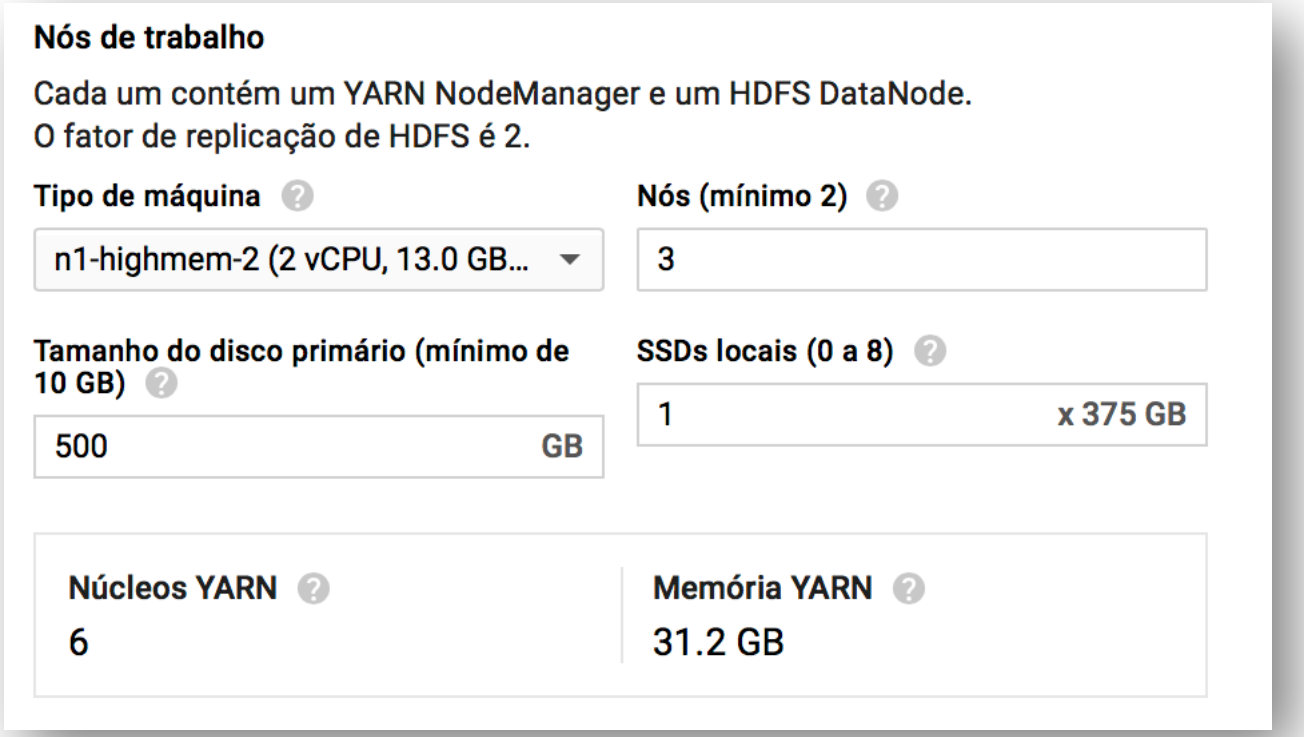

{ detalhes do cluster }

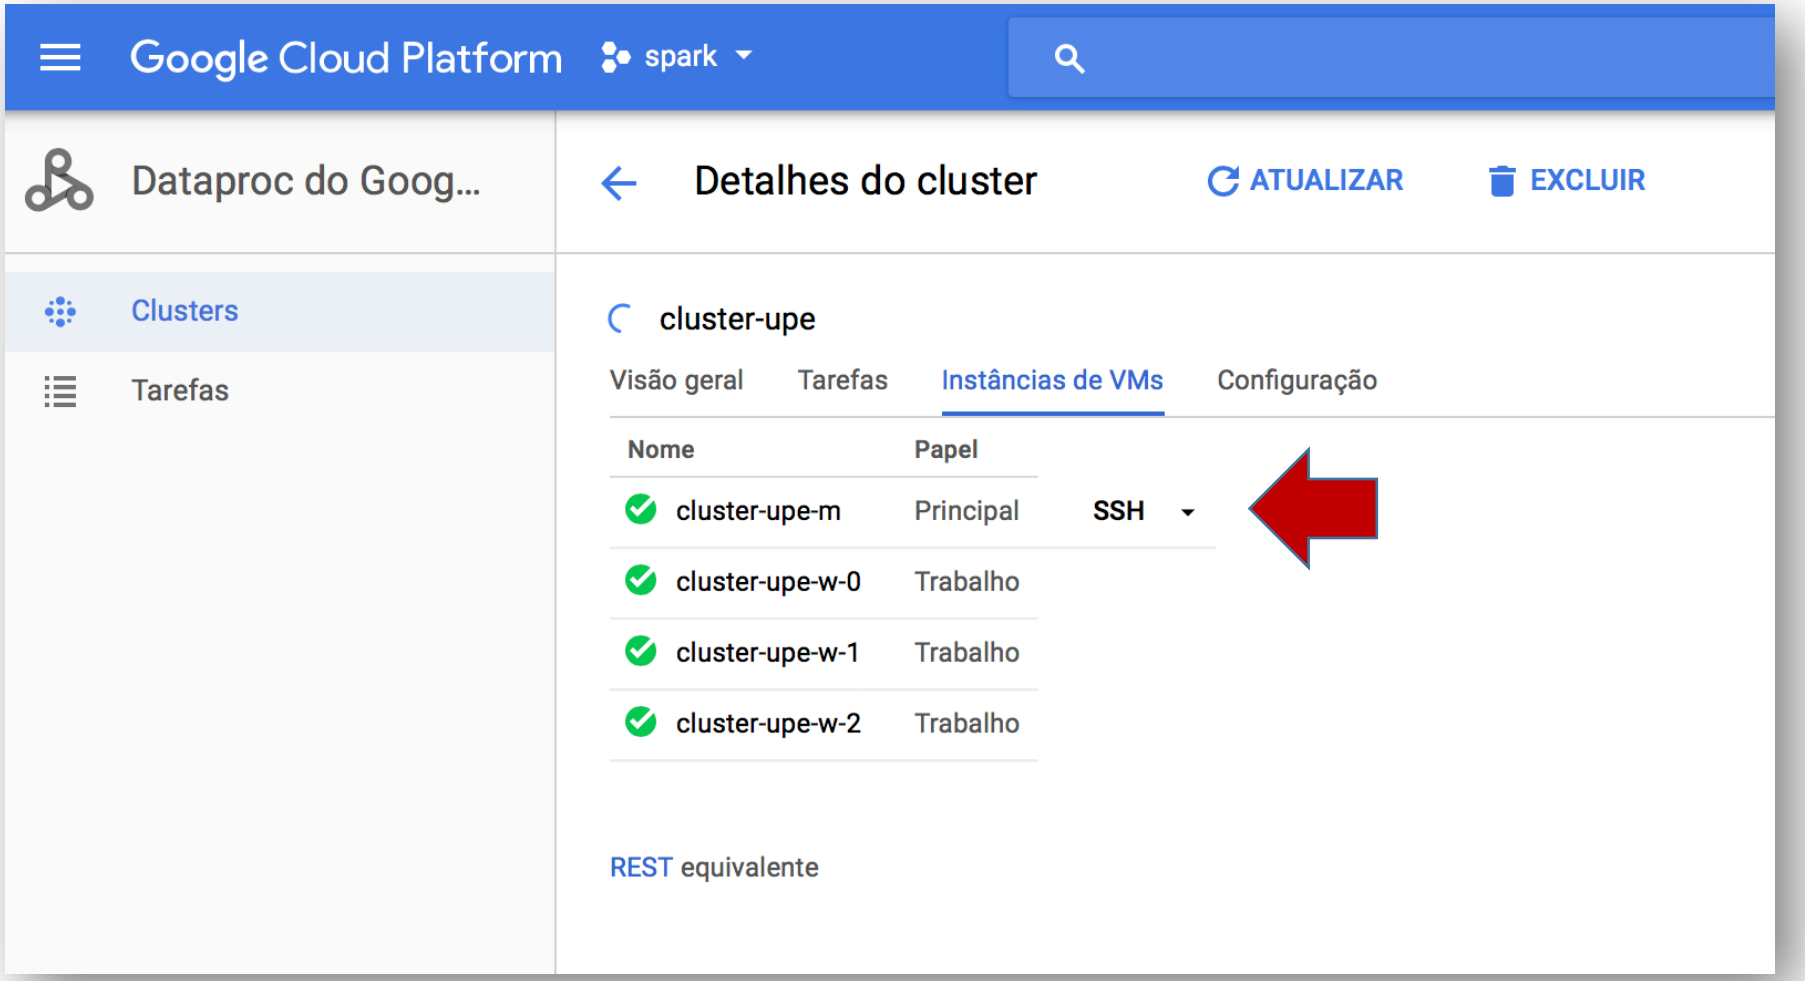

{ conexão SSH }

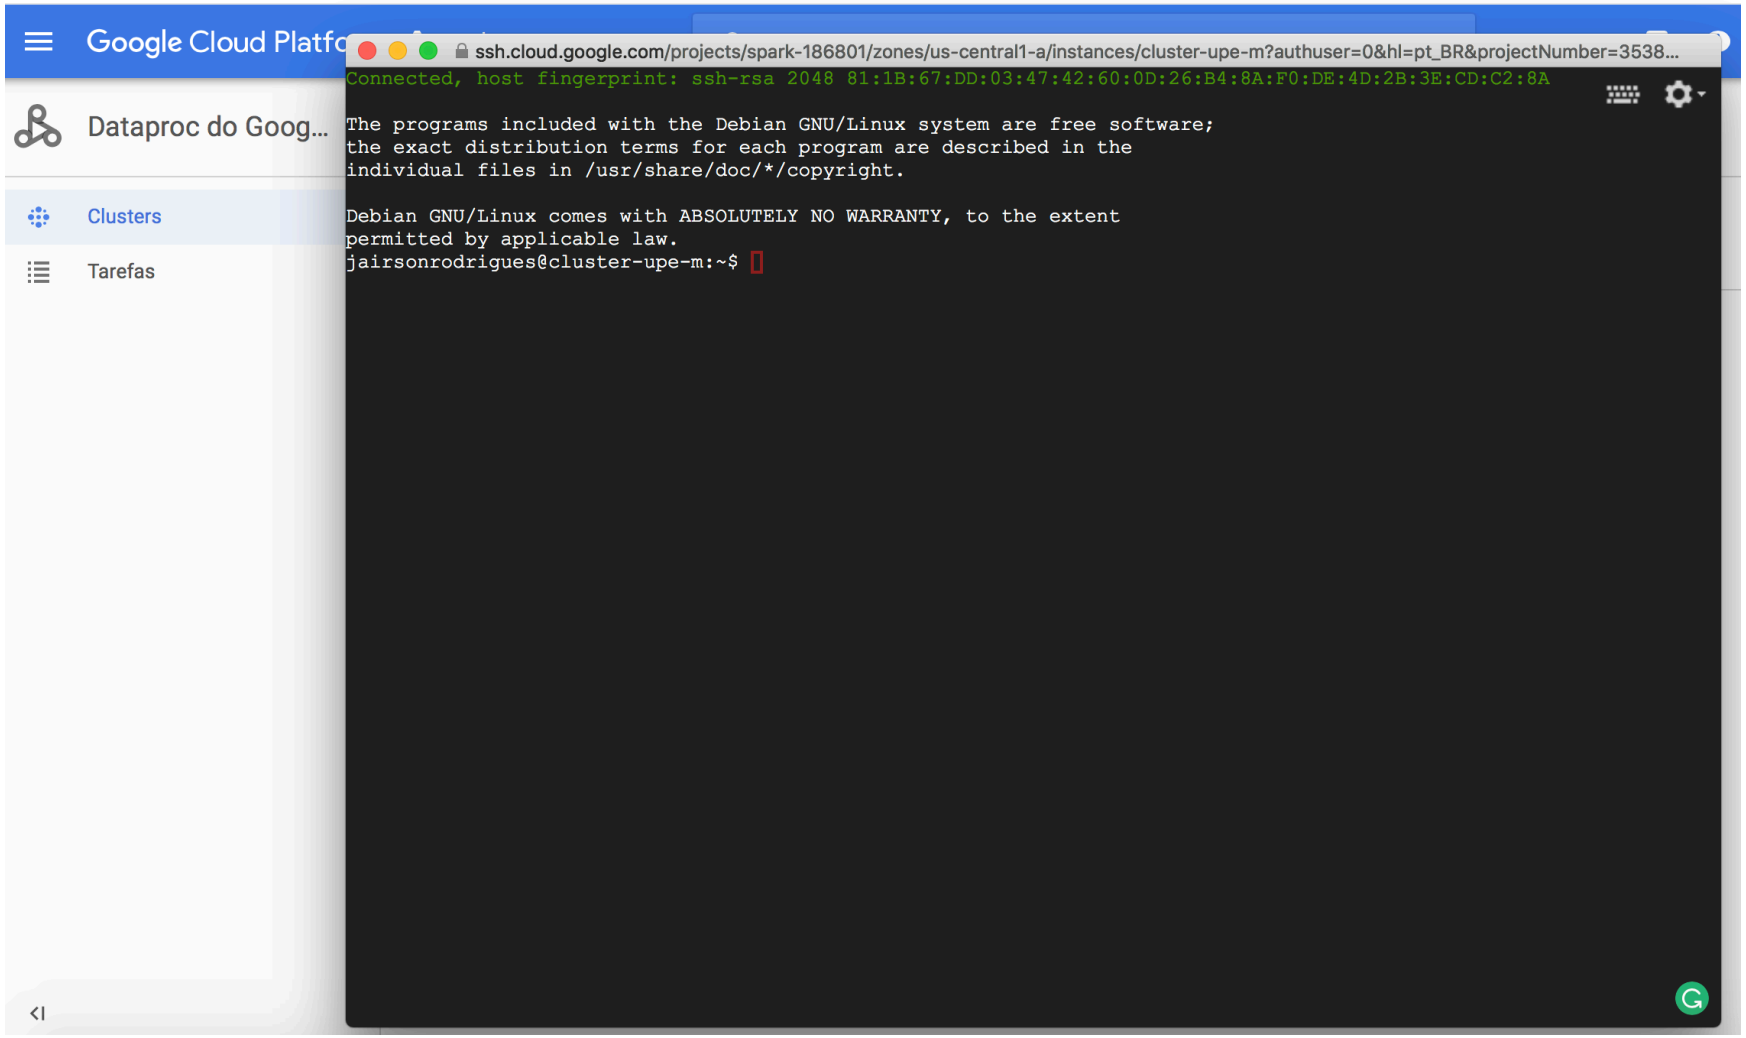
Google Cloud Platfo
Dataproc do Goog...
Clusters
Tarefas
ssh.cloud.google.com/projects/spark-186801/zones/us-central1-a/instances/cluster-upe-m?authuser=0&hl=pt_BR&projectNumber=3538...
Connected, host fingerprint: ssh-rsa 2048 81:1B:67:DD:03:47:42:60:0D:26:B4:8A:F0:DE:4D:2B:3E:CD:C2:8A
The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.
Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.
jairsonrodrigues@cluster-upe-m:~$

{ configurações iniciais }

- wget https://www.dropbox.com/s/r5xg2hi28g4s51f/kddcup_2.11-1.0.jar?dl=0
- wget https://github.com/SparkTC/spark-bench/releases/download/v55/spark-bench_2.1.1_0.2.2-RELEASE_55.tgz
- wget https://www.dropbox.com/s/0hi816m22uia2cw/spark-bench-env.sh?dl=0
- wget https://www.dropbox.com/s/98udl9vmi1fayzq/genkmeans.conf?dl=0
- wget https://www.dropbox.com/s/mwnad29s3zdtutx/kmeans.conf?dl=0
- wget http://kdd.ics.uci.edu/databases/kddcup99/kddcup.data.gz
- mv ~/kddcup_2.11-1.0.jar\?dl\=0 kddcup_2.11-1.0.jar
- mv ~/kmeans.conf?dl=0 ~/kmeans.conf
- mv ~/genkmeans.conf?dl=0 ~/genkmeans.conf
- gunzip kddcup.data.gz
- tar -xvf ~/spark-bench_2.1.1_0.2.2-RELEASE_55.tgz
- mv ~/spark-bench_2.1.1_0.2.2-RELEASE ~/spark-bench
- mv ~/spark-bench-env.sh\?dl\=0 ~/spark-bench/bin/spark-bench-env.sh
- mv kddcup.data kddcup.data.10
- hadoop fs -mkdir -p /kddcup/input/
- hadoop fs -copyFromLocal kddcup.data.10 /kddcup/input/

wget https://www.dropbox.com/s/hnfuinu3oose116/init.sh?dl=0
mv init.sh\?dl\=0 init.sh
chmod a+x init.sh; ./init.sh

{ geração de dados - kmeans -> 25 gb }

- ./spark-bench/bin/spark-bench.sh ~/genkmeans.conf

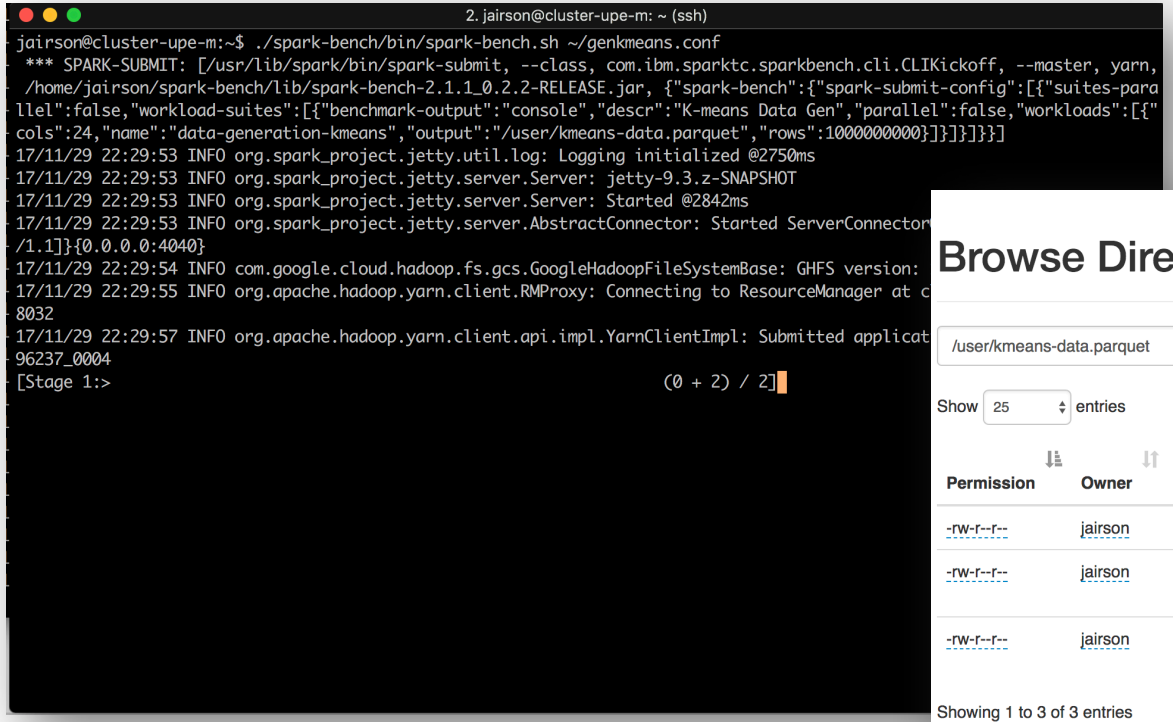

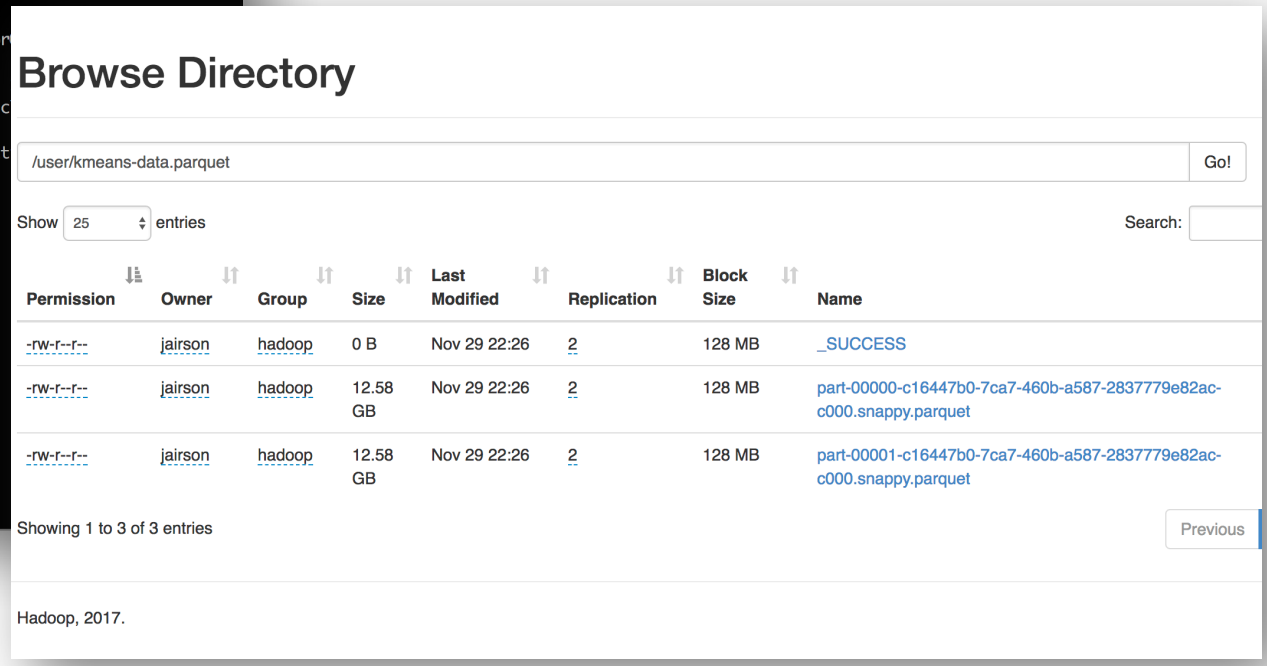

Proxy MAC OS: https://justpaste.it/1e30d

{ execução kmeans -> 8 gb }

- ./spark-bench/bin/spark-bench.sh ~/kmeans.conf

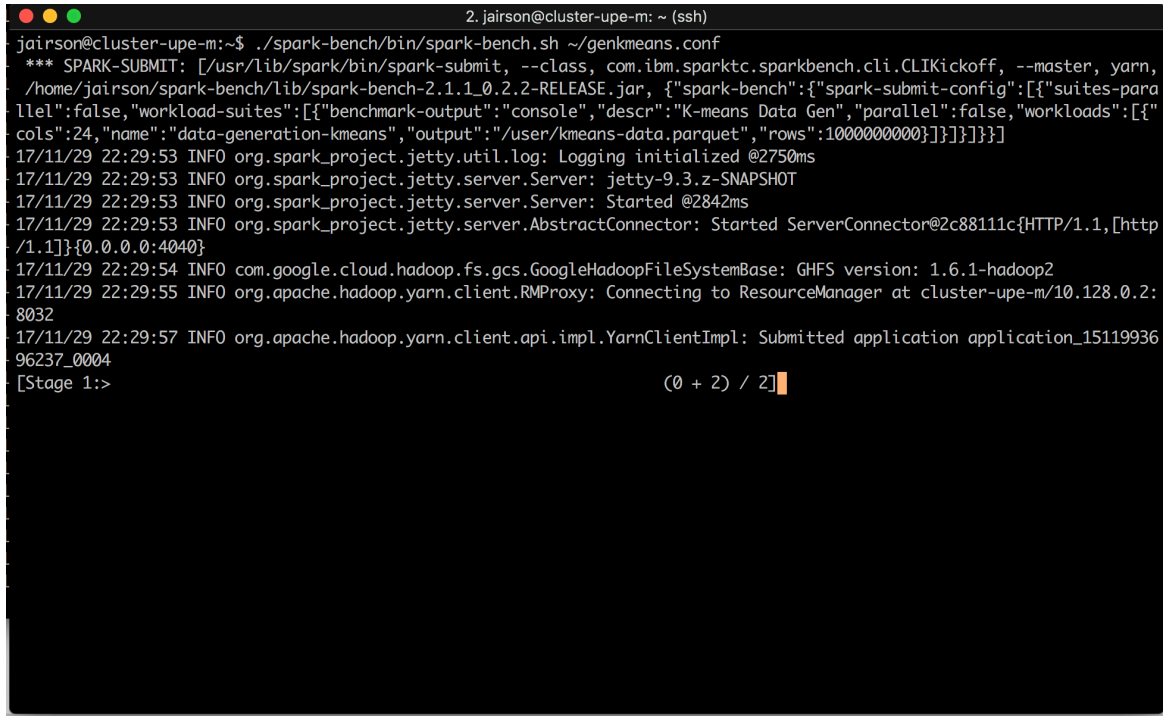

- 40 milhões de amostras
- 24 dimensões R^{24}
- # clusters: 5

{ envio de jobs / ETL / LogReg / Naïve Bayes }

- spark-submit --class "KDDCupETL" --master yarn --conf spark.serializer=org.apache.spark.serializer.KryoSerializer kddcup_2.11-1.0.jar
- spark-submit --class "KDDCupRL" --master yarn --conf spark.serializer=org.apache.spark.serializer.KryoSerializer kddcup_2.11-1.0.jar
- spark-submit --class "KDDCupNaiveBayes" --master yarn --conf spark.serializer=org.apache.spark.serializer.KryoSerializer kdkddcup_2.11-1.0.jar

{ envio de jobs - interface web }

- KDDCupETL
- KDDCupRL
- KDDCupNaiveBayes

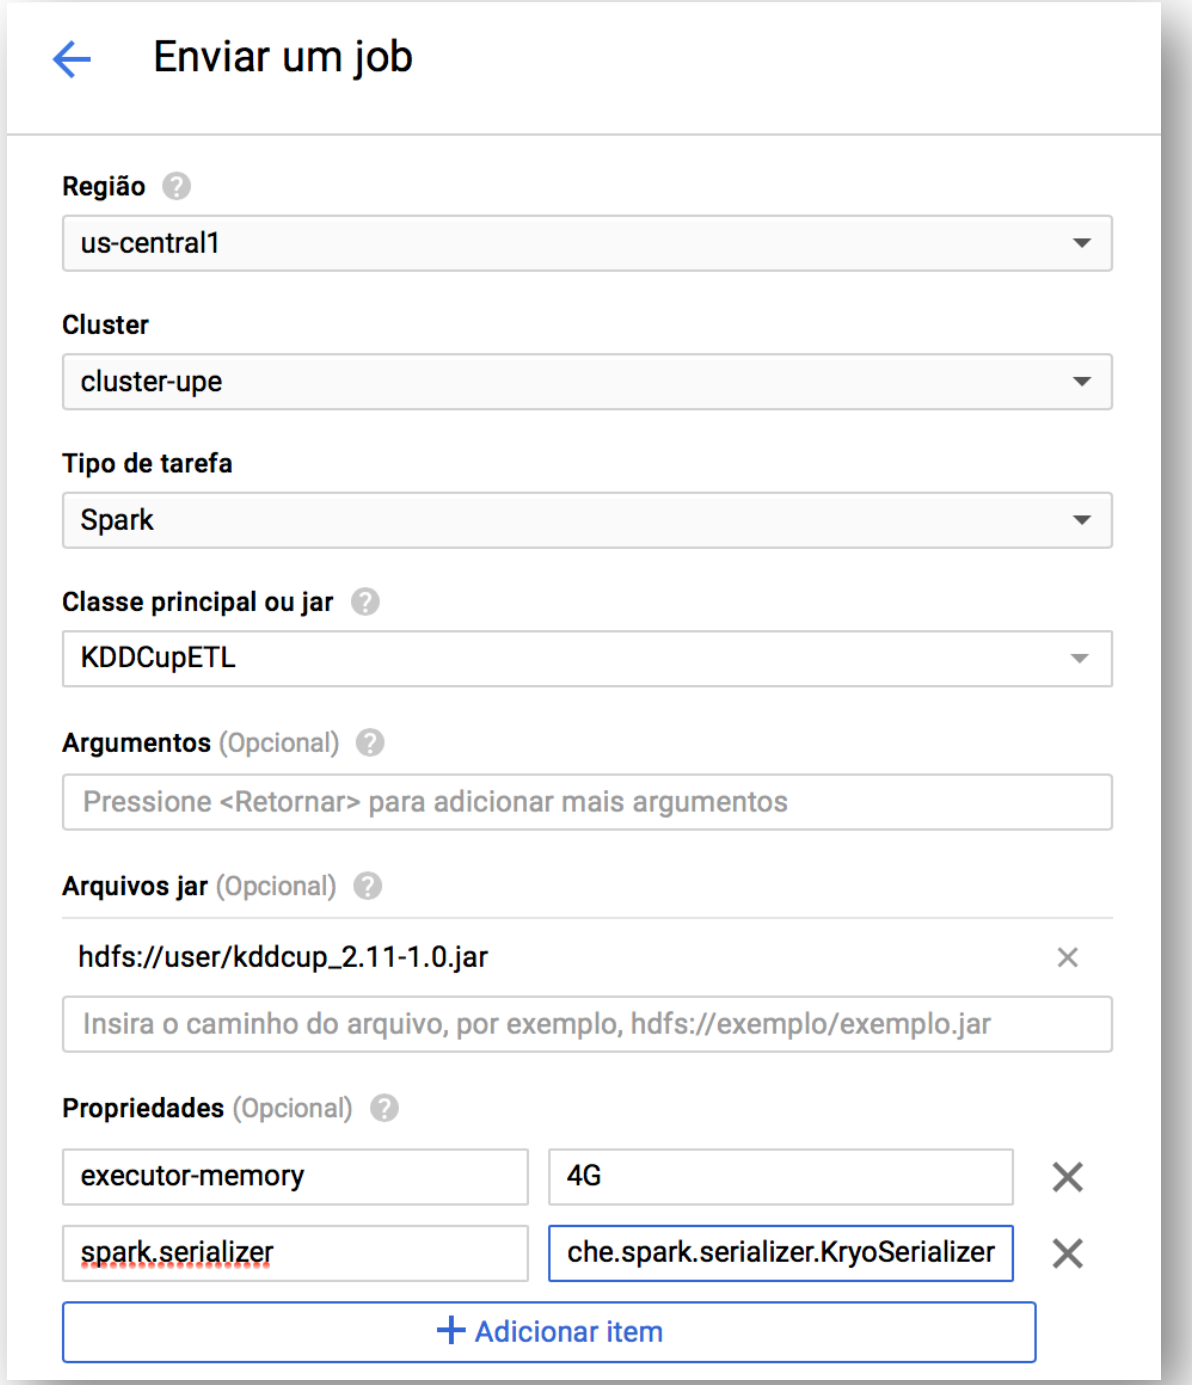

{ envio de jobs - interface web }